

# Parallel Programming

## Parallel Execution in CUDA (Part 2)

Phạm Trọng Nghĩa  
ptnghia@fit.hcmus.edu.vn

# Overview

---

**Apply knowledge of parallel execution in CUDA to write a fast CUDA program doing “reduction”**

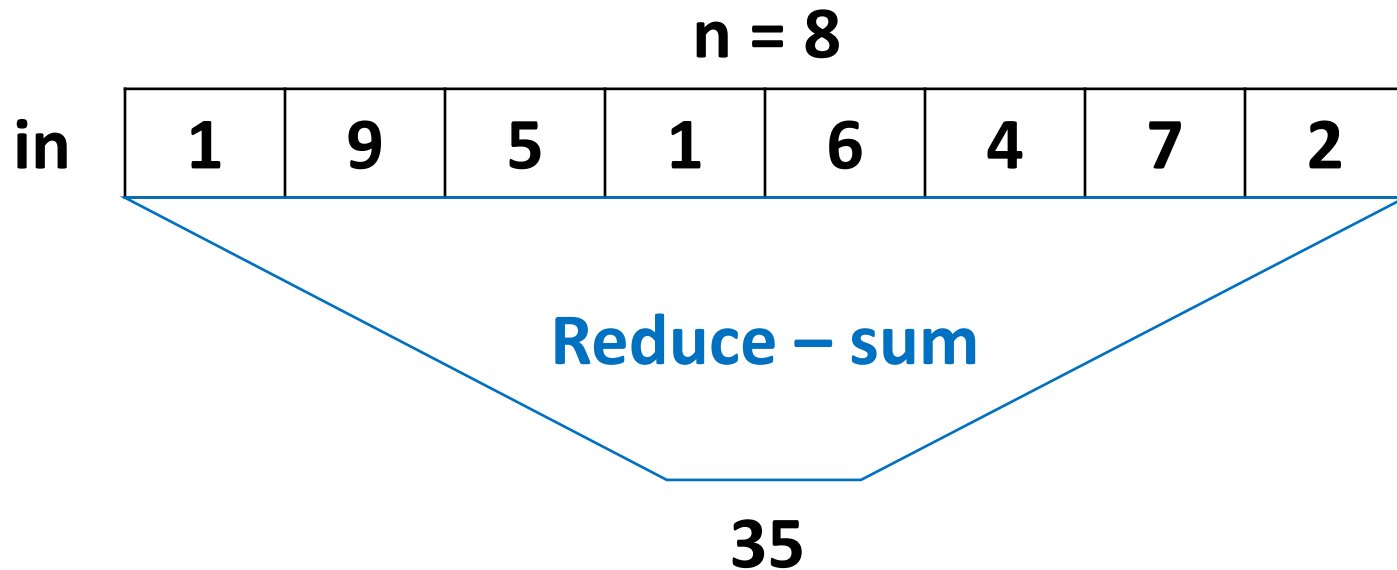
- The “reduction” task
- Sequential implementation
- Parallel implementation
  - Kernel function – 1<sup>st</sup> version
  - Kernel function – 2<sup>nd</sup> version: reduce warp divergence
  - Kernel function – 3<sup>rd</sup> version: reduce warp divergence + ...

# The “reduction” task

---

Input: an array **in** of **n** numbers

Output: sum (or product, max, min, ...) of these numbers



# Reduction algorithm

---

- Sum reduction sequential:

```
sum = 0;
for (int i = 0; i < N; i++){
    sum += input[i];
}
```

- General form of a reduction sequential:

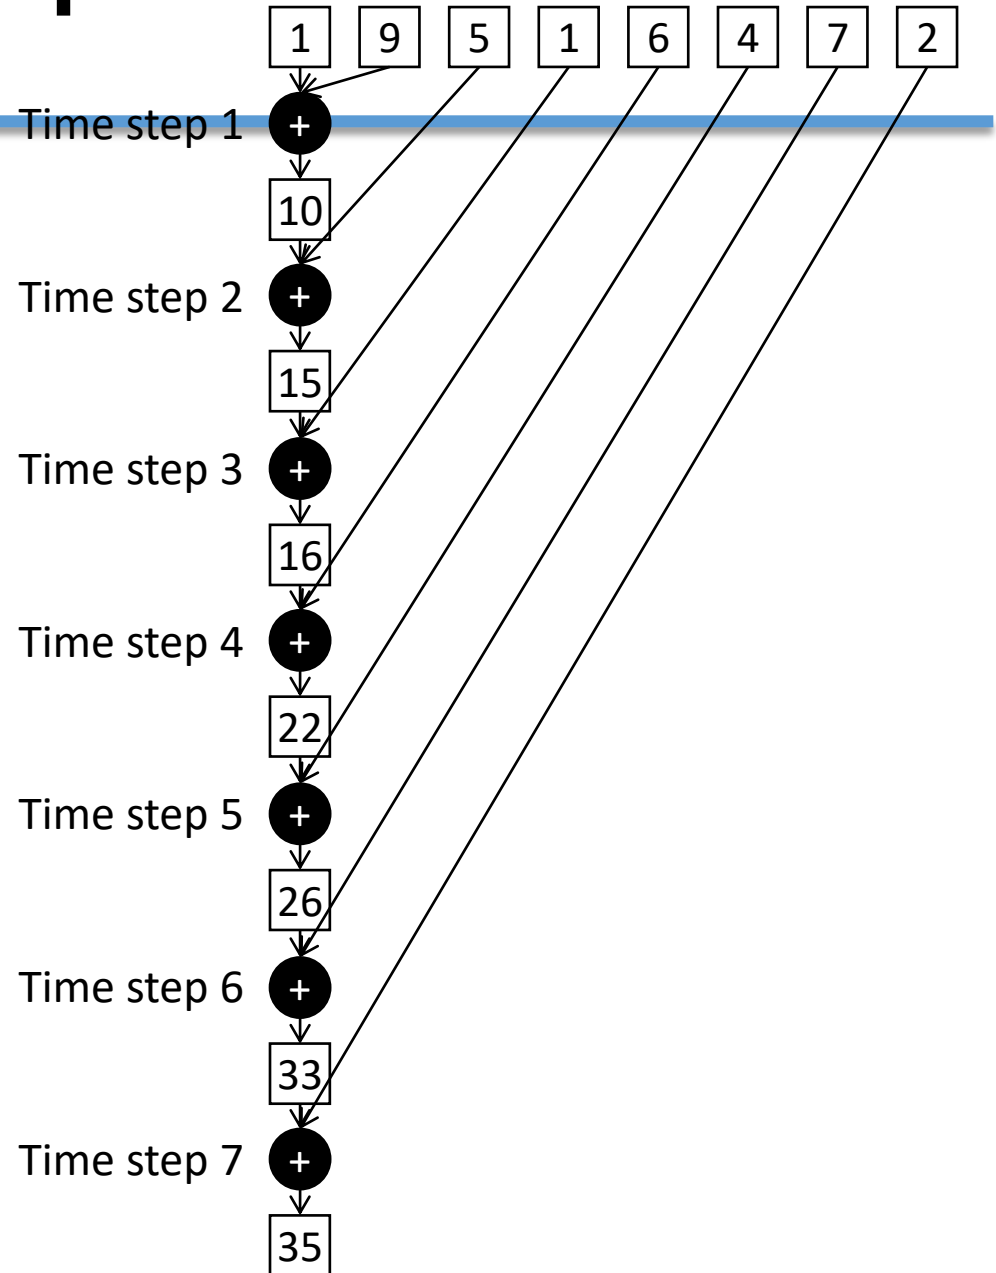
```
acc = IDENTITY;
for (int i = 0; i < N; i++){
    acc = Operator(acc, input[i]);
}
```

# Sequential implementation

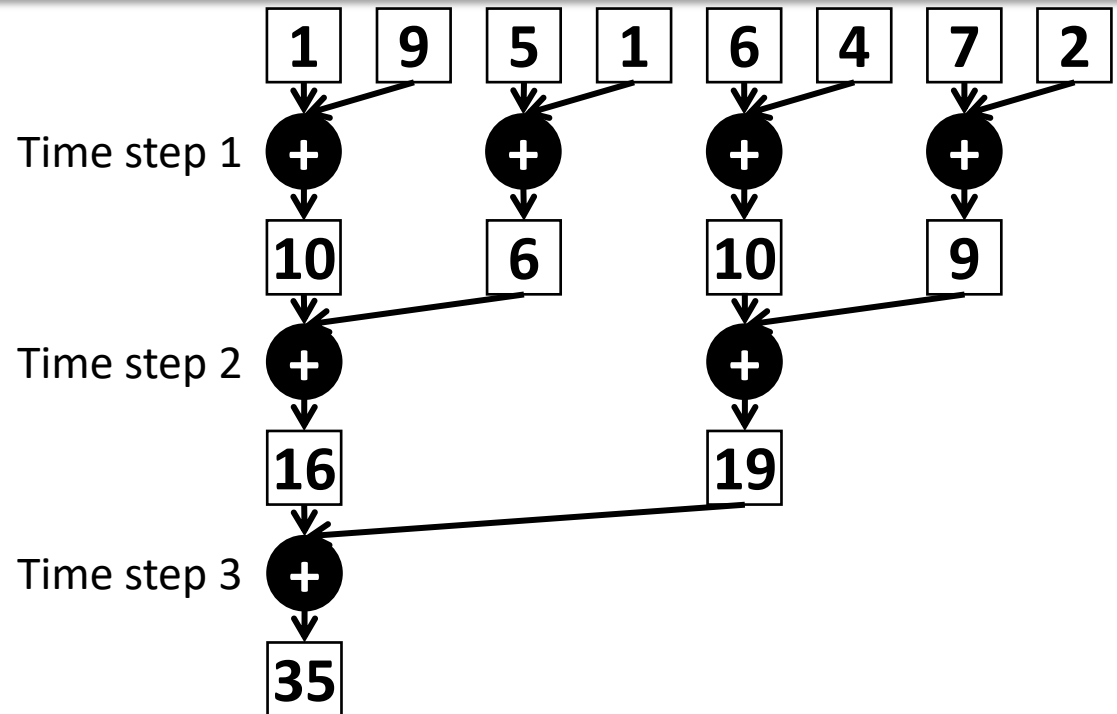
```
int reduceOnHost(int *in, int n)
{
    int s = in[0];
    for (int i = 1; i < n; i++)
        s += in[i];
    return s;
}
```

**Time (# time steps):**  $7 = n-1 = O(n)$

**Work (# pluses):**  $7 = n-1 = O(n)$



# Parallel implementation – idea



**Time: ?**

**Work: ?**

# Parallel implementation – idea

---

- For N input values, the reduction tree performs
  - $N/2 + N/4 + N/8 + \dots + 1 = N - 1$  operations
  - In  $\log(N)$  steps – 1,000,000 input values take 20 steps
    - Assuming that we have enough execution resources
  - Average Parallelism  $(N-1)/\log(N)$ 
    - For  $N = 1,000,000$ , average parallelism is 50,000
    - However, peak resource requirement is 500,000
    - This is not resource efficient
- This is a work-efficient parallel algorithm
  - The amount of work done is comparable to the an efficient sequential algorithm
  - Many parallel algorithms are not work efficient

# Reduction trees

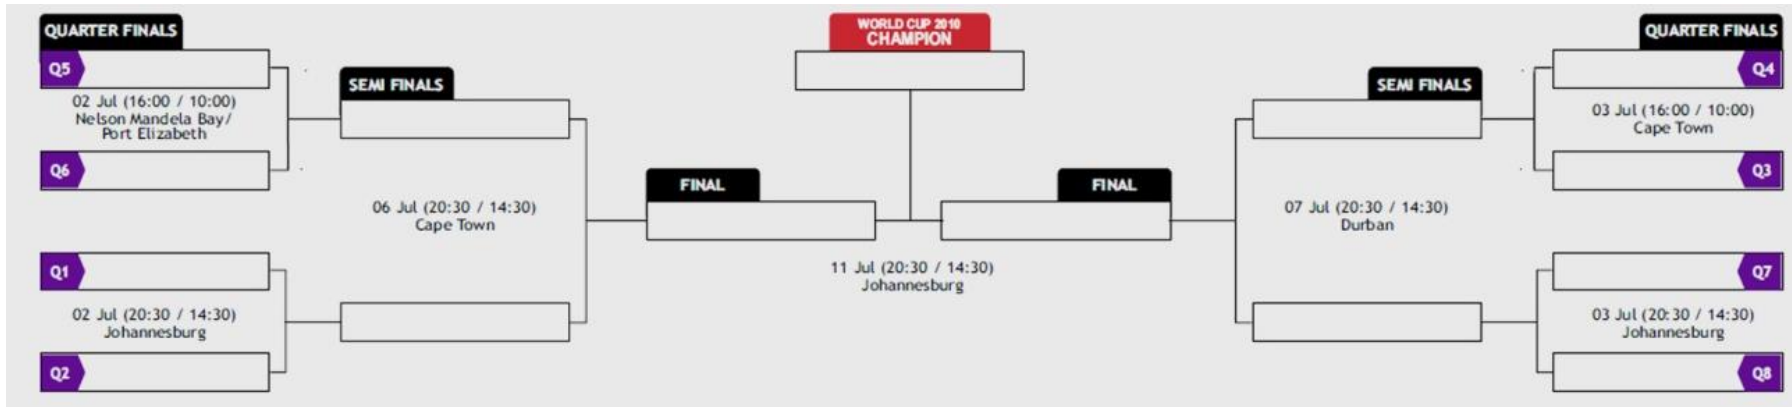
---

- Order of performing the operations will be changed (sequential → parallel)
  - Operator must be **associative**
  - Serial
    - ((((((3 max 1) max 7) max 0) max 4) max 1) max 6) max 3)
  - Paralell
    - ((3 max 1) max (7 max 0)) max ((4 max 1) max (6 max 3))
- We also need rearranges the order of the operands
  - Operator to be **commutative**



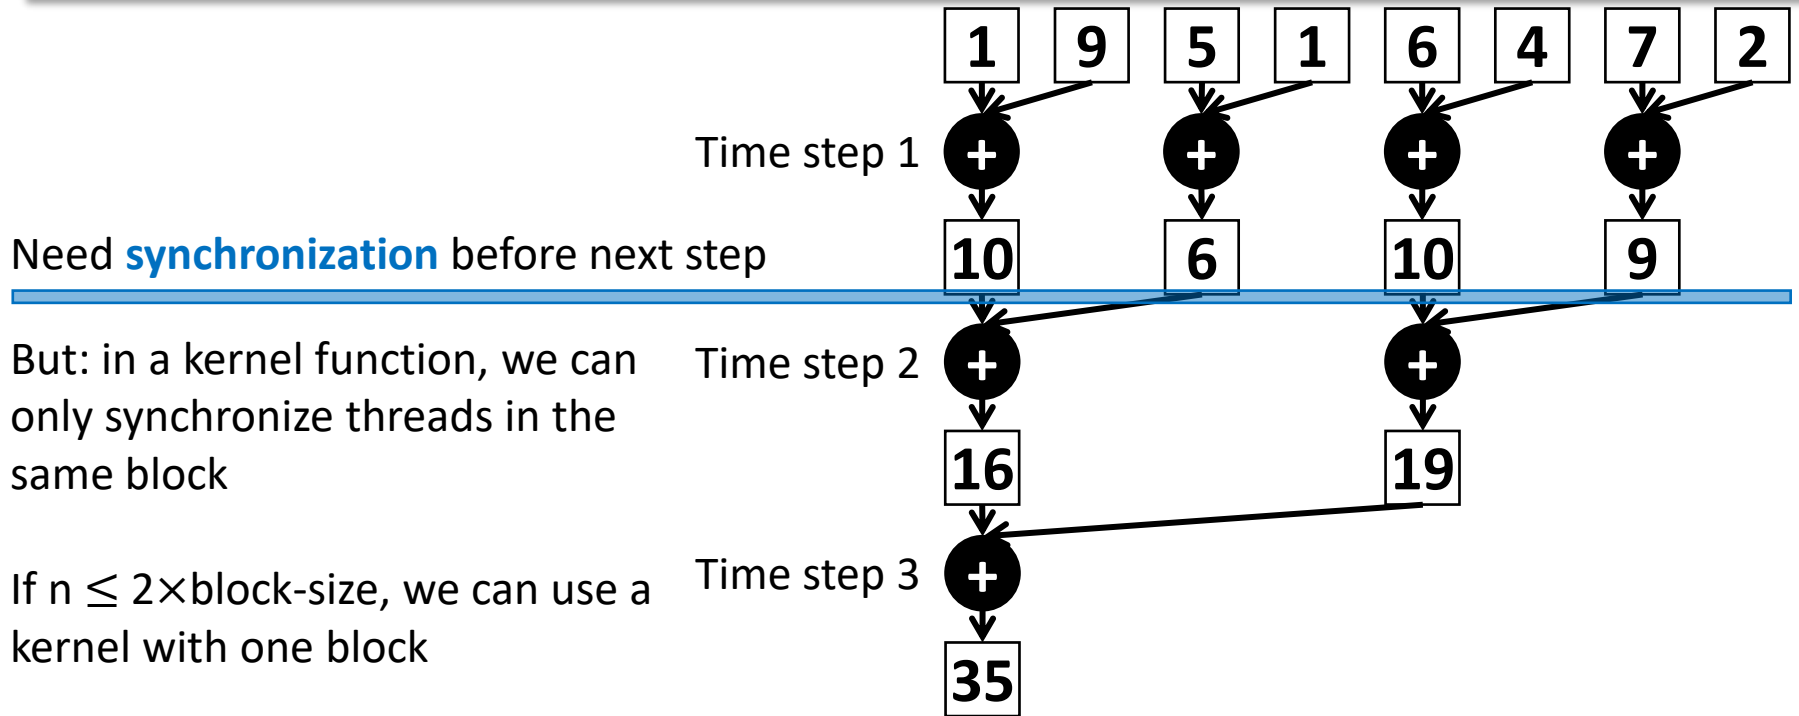
# Parallel Reduction in Real life

- Sports & Competitions: Max reduction



- Also use to process large input data sets (Google and Hadoop MapReduce frameworks)
  - There is no required order of processing elements in a data set (associative and commutative)
  - Partition the data set into smaller chunks
  - Have each thread to process a chunk
  - Use a reduction tree to summarize the results from each chunk into the final answer

# Parallel implementation – idea




**Time:**  $3 = \log_2 n = O(\log_2 n)$

**Work:**  $7 = n - 1 = O(n)$  = work of the sequential version

(Later, we will see tasks in which parallel implementations need to do more work than sequential)

# A simple reduction kernel

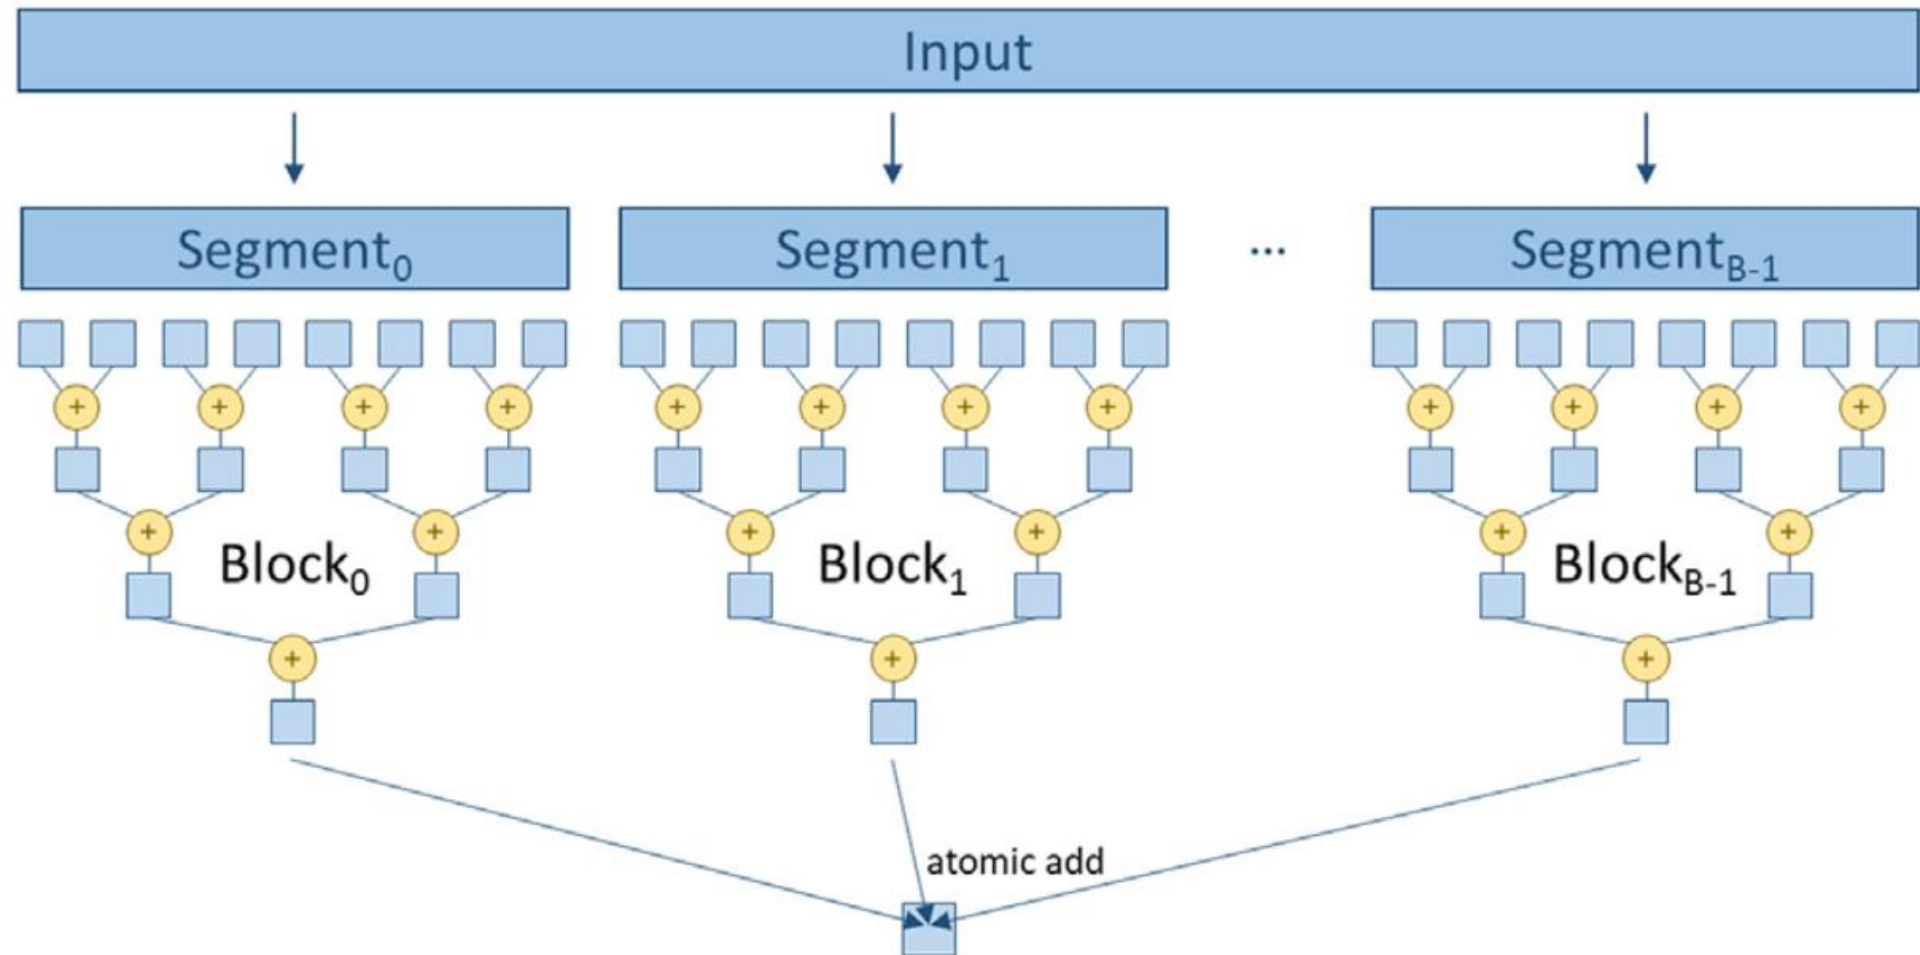


# A simple reduction kernel

---

```
__global__ void reduceBlksKernel0(int* in, int* out, int n) {
    int i = 2 * threadIdx.x;
    for (int stride = 1; stride <= blockDim.x; stride *= 2) {
        if (threadIdx.x % stride == 0)
            in[i] += in[i + stride];
        __syncthreads();
    }
    if (threadIdx.x == 0)
        *out = in[0];
}
```


# Hierarchical reduction for bigger input



# Parallel implementation

## – idea to reduce within each block

Consider a block of 4 threads



# Hierarchical reduction for arbitrary input length

```
__global__ void reduceBlksKernel1(int* in, int* out, int n){
    int i = blockIdx.x * 2 * blockDim.x + 2 * threadIdx.x;
    for (int stride = 1; stride <= blockDim.x; stride *= 2){
        if (threadIdx.x % stride == 0)
            if (i + stride < n)
                in[i] += in[i + stride];
        __syncthreads(); // Synchronize within each block
    }
    if (threadIdx.x == 0)
        atomicAdd(out, in[blockIdx.x * 2 * blockDim.x]);
}
```

Assume:  $2 \times \text{block-size} = 2^k$

# In each block, how many diverged warps? (not consider blocks in the edge)

---

- **Stride = 1:**

All threads are “on”

→ No diverged warp

- **Stride = 2:**

Only threads with `threadIdx.x % 2 == 0` are “on”

→ All warps are diverged

- **Stride = 4, 8, ..., 32:**

All warps are diverged

- **Stride = 64, 128, ....:**

# diverged warps decrease to 1




# Kernel function – 2<sup>nd</sup> version: reduce warp divergence

---

- **Idea:** reduce # diverged warps in each step by rearranging threads so that “on” threads are first adjacent threads
- **Example:** consider a block of 128 threads
  - Stride = 1: All 128 threads are “on”
  - Stride = 2: First 64 threads are “on”, the rest are “off”
  - Stride = 4: First 32 threads are “on”, the rest are “off”
  - ...

# Kernel function - 2<sup>nd</sup> version: reduce warp divergence

Consider a block of 4 threads



# Kernel function – 2<sup>nd</sup> version: reduce warp divergence

```
__global__ void reduceOnDevice2(int *in, int *out, int n)
{
    int numElemsBeforeBlk = blockIdx.x * blockDim.x * 2;

    for (int stride = 1; stride < 2 * blockDim.x; stride *= 2)
    {
        int i = numElemsBeforeBlk + ...;
        if (threadIdx.x ...)
            if (i + stride < n)
                in[i] += in[i + stride];
        __syncthreads(); // Synchronize within each block
    }

    if (threadIdx.x == 0)
        out[blockIdx.x] = in[numElemsBeforeBlk];
}
```

# Kernel function – 2<sup>nd</sup> version: reduce warp divergence

```
__global__ void reduceOnDevice2(int *in, int *out, int n)
{
    int numElemsBeforeBlk = blockIdx.x * blockDim.x * 2;

    for (int stride = 1; stride < 2 * blockDim.x; stride *= 2)
    {
        int i = numElemsBeforeBlk + threadIdx.x * 2 * stride;
        if (threadIdx.x < blockDim.x / stride)
            if (i + stride < n)
                in[i] += in[i + stride];
        __syncthreads(); // Synchronize within each block
    }


    if (threadIdx.x == 0)
        out[blockIdx.x] = in[numElemsBeforeBlk];
}
```

# Kernel function - 3<sup>rd</sup> version: reduce warp divergence + ?

Consider a block of 4 threads

Data of previous block

Data of next block



# **Kernel function - 3<sup>rd</sup> version: reduce warp divergence + ?**

Code: your next homework ;-)

# Reference

---

- [1] Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 2022
- [2] Cheng John, Max Grossman, and Ty McKercher. *Professional Cuda C Programming*. John Wiley & Sons, 2014



**THE END**