

Parallel Programming

Prefix sum (scan)

Phạm Trọng Nghĩa
ptnghia@fit.hcmus.edu.vn

Overview

- The “sort” task
- Sequential Radix Sort
- Parallel Radix Sort

The “sort” task

in

1	8	5	2	6	4	7	2
---	---	---	---	---	---	---	---

Stable sort

1	2	2	4	5	6	7	8
---	---	---	---	---	---	---	---

Unstable sort

1	2	2	4	5	6	7	8
---	---	---	---	---	---	---	---

We will focus on input array of **unsigned ints**

Overview

- The “sort” task
- Sequential Radix Sort
- Parallel Radix Sort

Sequential Radix Sort

Loop from bit b3 (least significant bit) to b1 (most significant bit):
Sort elements w.r.t. the current bit using a **stable** sort

	b1	b2	b3
1	0	0	1
0	0	0	0
5	1	0	1
2	0	1	0
6	1	1	0
4	1	0	0
7	1	1	1
2	0	1	0

	b1	b2	b3
0	0	0	0
2	0	1	0
6	1	1	0
4	1	0	0
2	0	1	0
1	0	0	1
5	1	0	1
7	1	1	1

	b1	b2	b3
0	0	0	0
4	1	0	0
1	0	0	1
5	1	0	1
2	0	1	0
6	1	1	0
2	0	1	0
7	1	1	1

	b1	b2	b3
0	0	0	0
1	0	0	1
2	0	1	0
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

DONE!

Sequential Radix Sort

- OK, Radix Sort works
- But is it efficient?

Yes, if we can make the stable sort in each loop efficient,
e.g. work = $O(n)$

- With unsigned int (32 bits),

Radix Sort's work $\approx 32n = O(n)$

- It's potentially even more efficient if we process $k > 1$ bits in each loop (and still keep the work in each loop at $O(n)$)

For simplicity, in this lecture, we just consider $k=1$ bit

Sort a binary array (corresponding to $k = 1$ bit in Radix Sort)

- Consider a binary input array:
binIn: 0 1 1 0 1 (n elements)
How to sort **stably** and **efficiently**?
- We will use Counting Sort
 - Compute the rank (the correct index in the output array) of each element (work = $O(n)$)
binIn: 0 1 1 0 1
ranks: 0 2 3 1 4

Rank of binIn[i] =
elements < binIn[i]
+ # elements before binIn[i] and = binIn[i]
 - Write each element to its rank in the output array (work = $O(n)$)

Sequential Radix Sort

Loop from Least Significant Bit to Most Significant Bit:

Sort elements w.r.t the current bit using Counting Sort (stably and efficiently)

Let's implement this ...

Overview

- The “sort” task
- Sequential Radix Sort
- **Parallel Radix Sort**

Sequential Radix Sort: parallelize?

Loop from Least Significant Bit to Most Significant Bit:

Sort elements w.r.t the current bit using Counting Sort

Parallelize✗

Parallelize✓

Sort a binary array using Counting Sort: parallelize?

- Consider a binary input array:
binIn: 0 1 1 0 1 (n elements)
How to sort **stably** and **efficiently**?
- We will use Counting Sort
 - Compute the rank (the correct index in the output array) of each element (work = $O(n)$)

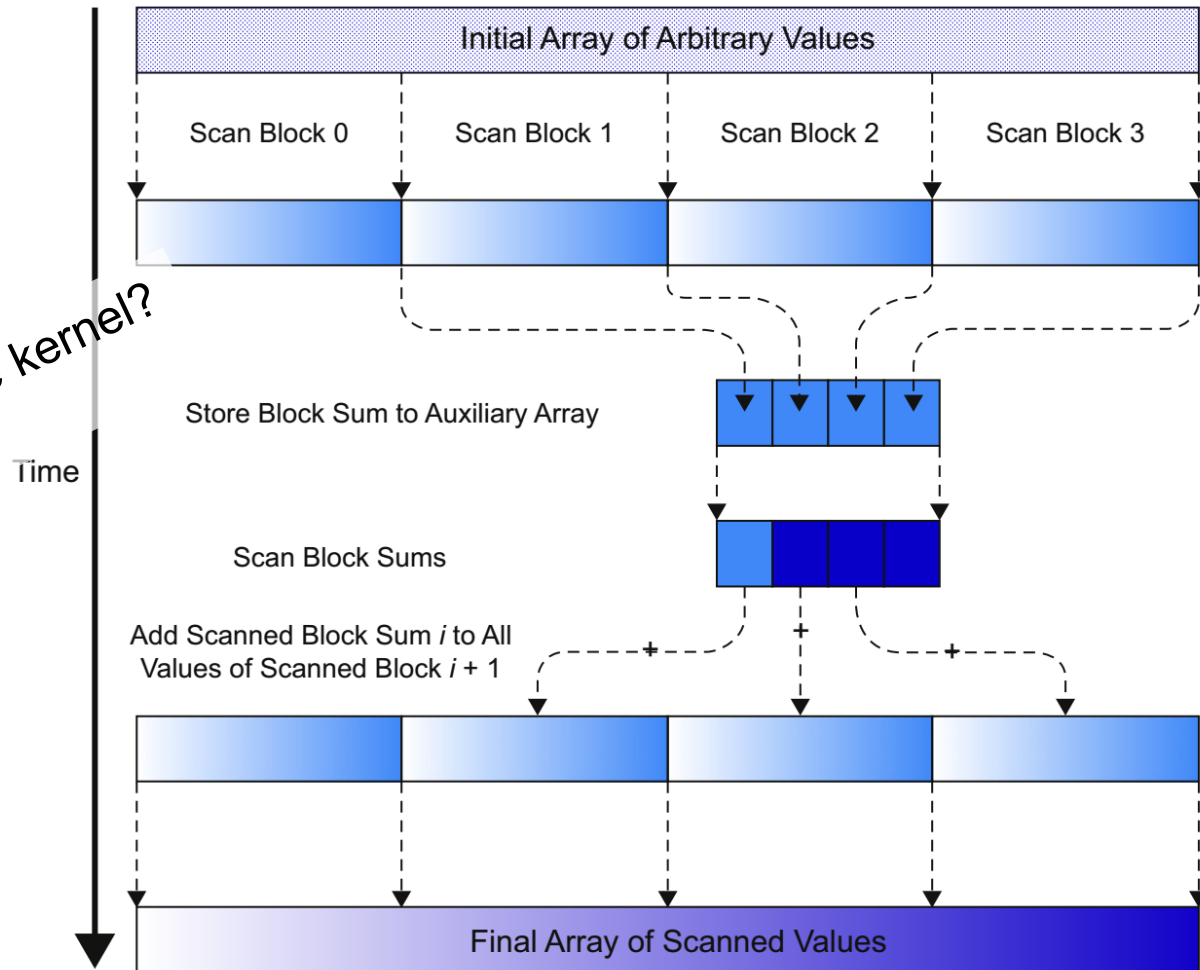
Parallelize ✓ • Compute # ones before each element:
binIn: 0 1 1 0 1
nOnesBefore: 0 0 1 2 2 Do exclusive scan

Parallelize ✓ • Compute rank:
if binIn[i] is 0: rank = i - nOnesBefore[i]
if binIn[i] is 1: rank = nZeros + nOnesBefore[i]
With nZeros = n - nOnesBefore[n-1] - binIn[n-1]
binIn: 0 1 1 0 1
ranks: 0 2 3 1 4

Parallelize ✓ • Write each element to its rank in the output array (work = $O(n)$)

Remember how do we implement scan in parallel?

Can we do all 3 phases in a single kernel?
Can we overlap these 3 phases?



Global scan in a single kernel

Block with index bi :

- Scan locally
- Wait until seeing the sign indicating block $bi-1$ has computed the sum of bi blocks ($0 \rightarrow bi-1$)

Get this sum, add this sum to block bi 's local sum, and turn on the sign indicating that block bi has computed the sum of $bi+1$ blocks ($0 \rightarrow bi$)

(Block $bi=0$ only needs to turn on the sign)

- Finish the rest of work: add the sum of bi blocks ($0 \rightarrow bi-1$) to block bi 's local scan

(Block $bi=0$ will not do this step)

Global scan in a single kernel

Block with index bi :

- Scan locally
- Wait until seeing the sign indicating block $bi-1$ has computed the sum of bi blocks ($0 \rightarrow bi-1$)

A possible situation:

Blocks $bi \rightarrow bi+N$ are assigned to available slots in SM, and wait for the result from block $bi-1$
Block $bi-1$ waits for an available slot in SM

→ Deadlock 😞

Solution: recompute block index bi , don't tie it with $blockIdx.x$
Finish the rest of bi and the sum of bi blocks ($0 \rightarrow bi-1$) to block bi 's local scan

(Block $bi=0$ will not do this step)

Global scan in a single kernel

Block with index bi :

- Get in-order block index bi
- Scan locally
- Wait until seeing the sign indicating block $bi-1$ has computed the sum of bi blocks ($0 \rightarrow bi-1$)

Get this sum, add this sum to block bi 's local sum, and turn on the sign indicating that block bi has computed the sum of $bi+1$ blocks ($0 \rightarrow bi$)

(Block $bi=0$ only needs to turn on the sign)

- Finish the rest of work: add the sum of bi blocks ($0 \rightarrow bi-1$) to block bi 's local scan

(Block $bi=0$ will not do this step)

Get in-order block index bi

- *blkCount1*: dùng để gán block ID (bid) mới. Giá trị đầu = 0
 - Block đầu tiên chạy sẽ có giá trị (bid = 0), *blkCount1++*
 - Block thứ 2 sẽ có bid = 1, *blkCount1++*
 - ...
- Chỉ thread 0 cần tính phần này vào biến share. Sau đó các thread khác lấy giá trị share này vào biến register của mình.

```
__device__ int blkCount1 = 0;
__device__ int blkCount2 = 0;
//...
if (threadIdx.x == 0)
    s_data[0] = atomicAdd(&blkCount1,1);
__syncthreads();
int bid = s_data[0];
```

Get in-order block index bi

```
__device__ int blkCount1 = 0;
__device__ int blkCount2 = 0;
//...
if (threadIdx.x == 0){
    blkSums[bid] = s_data[2 * blockDim.x - 1];
    if (bid > 0){
        while (atomicAdd(&blkCount2, 0) < bid) {}
        s_data[blockDim.x * 2] = blkSums[bid - 1];
        blkSums[bid] += s_data[blockDim.x * 2];
        __threadfence();
    }
    atomicAdd(&blkCount2, 1);
}
__syncthreads();
```

Read more:

- [Document about __threadfence](#)
- [Document about volatile](#)

Global scan in a single kernel

Block with index bi :

- Serialize between blocks*
- Get in-order block index bi
 - Scan locally
 - Wait until seeing the sign indicating block $bi-1$ has computed the sum of bi blocks ($0 \rightarrow bi-1$)
Get this sum, add this sum to block bi 's local sum, and turn on the sign indicating that block bi has computed the sum of $bi+1$ blocks ($0 \rightarrow bi$)
(Block $bi=0$ only needs to turn on the sign)
 - Finish the rest of work: add the sum of bi blocks ($0 \rightarrow bi-1$) to block bi 's local scan
(Block $bi=0$ will not do this step)

Inclusive scan $\overset{?}{\rightarrow}$ exclusive scan

Implement parallel Radix Sort using global scan in a single kernel

The upcoming HW4 ;-)

Radix Sort for signed ints

- Sign bit is MSB (Most Significant Bit)
 - MSB = 0: positive number
Signed int = unsigned int
 - MSB = 1: negative number
Signed int = unsigned int - $2^{\text{\#bits-of-signed-int}}$
- If we use Radix Sort for unsigned ints, it'll be wrong
- One solution:
 - Convert signed ints to unsigned ints
 - Run Radix Sort for unsigned ints
 - Convert results back to signed ints

Radix Sort for floats

- Need to understand how floats are represented
- Idea is similar to signed ints:
 - Convert floats to unsigned ints
 - Run Radix Sort for unsigned ints
 - Convert results back to floats

Reference

- [1] Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 2022
- [2] Cheng John, Max Grossman, and Ty McKercher. *Professional Cuda C Programming*. John Wiley & Sons, 2014
- [3] Illinois GPU course

<https://wiki.illinois.edu/wiki/display/ECE408/ECE408+Home>



THE END